

Ankit Kumar

+91-9667267972 | ankit.k.j1999@gmail.com

Gurugram, Haryana, India | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | [DataPathSala](#)

PROFILE

Data Engineer with 3+ years of experience building automated ETL pipelines, streaming systems, and cloud data infrastructure across **hospitality, e-commerce, and SaaS**. Skilled in **Python, SQL, PySpark, Scala, AWS, Terraform, and GitHub Actions**, with experience across Redshift, S3, MongoDB, PostgreSQL, and AWS Glue.

TECHNICAL SKILLS

Programming Languages: Python, SQL, Scala, Bash

Big Data & Processing: PySpark, Apache Spark, Spark SQL, GraphX, Pandas, Databricks

Cloud (AWS): S3, Redshift, Glue, Kinesis, Lambda, SNS, CloudWatch, EC2, ECR

Streaming & Search: AWS Kinesis, RabbitMQ, Elasticsearch / OpenSearch

Databases: PostgreSQL, MongoDB, Redis, MySQL

Orchestration & CI/CD: Apache Airflow, GitHub Actions, Terraform

DevOps & Backend: Docker, Linux, Node.js, Express, FastAPI

EXPERIENCE

Lemon Tree Hotels

July 2026 - Present

Data Engineer

India

- Implemented **Terraform + GitHub Actions CI/CD** for AWS Lambda deployments, standardizing source control, infrastructure provisioning, and repeatable releases through an organized Git repository.
- Migrated approximately **65 AWS Glue jobs** from ad-hoc AWS configurations into a systematic, pipeline-oriented Git repository; developed **300+ unit and integration tests** to improve deployment confidence and maintainability.
- Rewrote the customer deduplication pipeline, removing unnecessary processing and legacy logic; migrated the graph-based implementation to **PySpark** and implemented **Breadth-First Search (BFS)** for scalable traversal.
- Resolved intermittent Tableau data-source refresh failures caused by CPU contention around **Tableau Bridge** and a Redshift warehouse in a private VPC; changed concurrent refreshes to controlled sequential execution, improving reliability.

Zalon

April 2024 - July 2026

Data Engineer

Gurugram, India

- Led a team of **4 Data Engineers** in designing, developing, and optimizing data pipelines and ETL processes supporting an e-commerce and SaaS platform.
- Built scalable real-time event ingestion using **AWS Kinesis and Lambda**, staging application events in Amazon S3 and processing **5M+ daily records** with zero data loss.
- Architected **10+ AWS Glue (PySpark) and Python ETL pipelines** covering **500,000+ SKUs**; migrated legacy MongoDB records into centralized Redshift/PostgreSQL warehouses with schema validation and deduplication.

Techtonic

September 2023 - April 2024

Data Engineer

Ghaziabad, India

- Built Python-based REST API ingestion pipelines into **PostgreSQL**, processing **50K+ records daily** with validation, deduplication, and SQL reporting views for an enterprise SaaS platform.
- Leveraged Firebase Analytics event data to identify user drop-offs, contributing to a **21% improvement in retention**; archived processed exports to AWS S3 for audit compliance and versioned backup.

Tekniko

June 2023 - September 2023

Data Engineering Intern

Noida, India

- Built Python and Pandas pipelines to clean, deduplicate, and schema-normalize **100K+ records** of raw activity data, automating PostgreSQL loading and eliminating **10+ hours/week** of manual reporting effort.

EDUCATION

Lovely Professional University

Punjab, India

B.Tech. in Computer Science & Engineering

2019 - 2023

CERTIFICATIONS & ACCOMPLISHMENTS

AWS Certified Cloud Practitioner

Jan 2025 - Jan 2028

Published **65+ technical articles** on GeeksforGeeks, accumulating **7M+ total views**.

Rank #4 on DataPathSala (Data Engineering platform) — **301/439 problems solved, 64-day streak, Level 8 Grandmaster**.