

Ankit Kumar

+91-9667267972 | ankit.k.j1999@gmail.com
Gurugram, Haryana, India | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFILE

Data Engineer with 3+ years of experience building automated ETL pipelines, streaming systems, and cloud data infrastructure across **hospitality, e-commerce, and SaaS**. Skilled in **Python, SQL, PySpark, Scala, AWS, Terraform, and GitHub Actions**, with experience across Redshift, S3, MongoDB, PostgreSQL, and AWS Glue.

TECHNICAL SKILLS

Programming Languages: Python, SQL, Scala, Bash
Big Data & Processing: PySpark, Apache Spark, Spark SQL, GraphX, Pandas, Databricks
Cloud (AWS): S3, Redshift, Glue, Kinesis, Lambda, SNS, CloudWatch, EC2, ECR
Streaming & Search: AWS Kinesis, RabbitMQ, Elasticsearch / OpenSearch
Databases: PostgreSQL, MongoDB, Redis, MySQL
Orchestration & CI/CD: Apache Airflow, GitHub Actions, Terraform
DevOps & Backend: Docker, Linux, Node.js, Express, FastAPI

EXPERIENCE

- Lemon Tree Hotels** July 2026 - Present
Data Engineer India
 - Implemented **Terraform + GitHub Actions CI/CD** for AWS Lambda deployments, standardizing source control, infrastructure provisioning, and repeatable releases through an organized Git repository.
 - Migrated approximately **65 AWS Glue jobs** from ad-hoc AWS configurations into a systematic, pipeline-oriented Git repository; developed **300+ unit and integration tests** to improve deployment confidence and maintainability.
 - Rewrote the customer deduplication pipeline, removing unnecessary processing and legacy logic; migrated the graph-based implementation to **PySpark** and implemented **Breadth-First Search (BFS)** for scalable traversal.
 - Resolved intermittent Tableau data-source refresh failures caused by CPU contention around **Tableau Bridge** and a Redshift warehouse in a private VPC; changed concurrent refreshes to controlled sequential execution, improving reliability.
- Zalon** April 2024 - July 2026
Data Engineer Gurugram, India
 - Led a team of **4 Data Engineers** in designing, developing, and optimizing data pipelines and ETL processes supporting an e-commerce and SaaS platform.
 - Built scalable real-time event ingestion using **AWS Kinesis and Lambda**, staging application events in Amazon S3 and processing **5M+ daily records** with zero data loss.
 - Architected **10+ AWS Glue (PySpark) and Python ETL pipelines** covering **500,000+ SKUs**; migrated legacy MongoDB records into centralized Redshift/PostgreSQL warehouses with schema validation and deduplication.
- Techtonic** September 2023 - April 2024
Data Engineer Ghaziabad, India
 - Built Python-based REST API ingestion pipelines into **PostgreSQL**, processing **50K+ records daily** with validation, deduplication, and SQL reporting views for an enterprise SaaS platform.
 - Leveraged Firebase Analytics event data to identify user drop-offs, contributing to a **21% improvement in retention**; archived processed exports to AWS S3 for audit compliance and versioned backup.
- Tekniko** June 2023 - September 2023
Data Engineering Intern Noida, India
 - Built Python and Pandas pipelines to clean, deduplicate, and schema-normalize **100K+ records** of raw activity data, automating PostgreSQL loading and eliminating **10+ hours/week** of manual reporting effort.

EDUCATION

- Lovely Professional University** Punjab, India
B.Tech. in Computer Science & Engineering 2019 - 2023

CERTIFICATIONS & ACCOMPLISHMENTS

- AWS Certified Cloud Practitioner** Jan 2025 - Jan 2028
Published **65+ technical articles** on GeeksforGeeks, accumulating **7M+ total views**.